

```
import torch
torch.cuda.is_available()
```

→ True

```
import os
import re
import json
import random
import string
from dataclasses import dataclass
from typing import Dict, List, Union, Optional
```

```
import torch
import torchaudio
import librosa
import evaluate
from datasets import load_dataset, Audio, DatasetDict
```

```
# Audio parameters
TARGET_SAMPLING_RATE = 16000
```

```
# Training output dir
OUTPUT_DIR = "Ed-168/wav2vec2-large-xls-r-300m-hi"
```

```
# Training hyperparameters (tune for your budget)
BATCH_SIZE = 1
GRAD_ACCUM = 16
LEARNING_RATE = 3e-5 # Good starting LR for ASR, tune lower if model is unstable
NUM_TRAIN_EPOCHS = 40
EVAL_STRATEGY = "steps"
EVAL_STEPS = 1000 # Evaluate less frequently to save memory
SAVE_STEPS = 1000 # Save less frequently to reduce disk I/O
LOGGING_STEPS = 50
WARMUP_RATIO = 0.05
FP16 = torch.cuda.is_available() # Enable mixed precision
```

```
# # If you want to push to the Hub, set these:
# PUSH_TO_HUB = True
# HF_REPO_ID = "Ed-168/Fine-tuned-wav2vec2-BERT-indian-languages" # e.g. "username/w2vbert-hi-ctc-cv17"
```

```
from datasets import load_dataset
```

```
common_voice_train = load_dataset("mozilla-foundation/common_voice_16_0", "hi", split="train+validation")
common_voice_test = load_dataset("mozilla-foundation/common_voice_16_0", "hi", split="test")
print(common_voice_train)
print(common_voice_test)
```

→ Dataset({
 features: ['client_id', 'path', 'audio', 'sentence', 'up_votes', 'down_votes', 'age', 'gender', 'accent', 'locale', 'segment',
 num_rows: 7084
})
Dataset({
 features: ['client_id', 'path', 'audio', 'sentence', 'up_votes', 'down_votes', 'age', 'gender', 'accent', 'locale', 'segment',
 num_rows: 3107
})

```
len(common_voice_train)
```

→ 7084

```
NUM_TRAIN_SAMPLES = 1500
NUM_TEST_SAMPLES = 750
```

```
common_voice_train = common_voice_train.select(range(NUM_TRAIN_SAMPLES))
common_voice_test = common_voice_test.select(range(NUM_TEST_SAMPLES))
```

```
common_voice_train = common_voice_train.remove_columns(["accent", "age", "client_id", "down_votes", "gender", "locale", "segment", "up_votes"])
common_voice_test = common_voice_test.remove_columns(["accent", "age", "client_id", "down_votes", "gender", "locale", "segment", "up_votes"])
```

```
def display_samples(data):
    for i in range(10):
```

```
print(f"{i+1} {data[i]['sentence']}")
display_samples(common_voice_train)
```

- 1 हमने उसका जन्मदिन मनाया।
- 2 साउथ दिल्ली नगर निगम सख्त, शॉपिंग मॉल के बाहर नहीं दिखेंगे होर्डिंग
- 3 उत्तर कोरिया ने अमेरिका को दो हमले की धमकी
- 4 अगले कमरे में अनेक रोमन मूर्तियाँ हैं।
- 5 तुम ने टॉम को कहाँ भेज दिया?
- 6 सदियों के आने से दिन छोटे होते जाते हैं।
- 7 मुझे और वक्त दो।
- 8 कंगना के वकील ने कहा, पुलिस ने किसी लैपटॉप की डिमांड नहीं की है
- 9 क्या सवाल है!
- 10 वह अच्छा राजा था।

```
import re
```

```
chars_to_ignore_regex = r"[\"'\\(\)\[\]\{\}\<|>|-|-|-|-|-|-|. , |?|!|:|;|\\|d|@|#|$|%|^&*|+|= _\\\\\\/|~` ]+" 
```

```
def normalize_text(batch):
    text = batch["sentence"]
    text = text.lower()
    text = re.sub(chars_to_ignore_regex, " ", text)
    text = re.sub(r"\\s+", " ", text).strip()
    batch["sentence"] = text
    return batch

common_voice_train = common_voice_train.map(normalize_text)
common_voice_test = common_voice_test.map(normalize_text)
display_samples(common_voice_train)
```

```

Map: 0% | 0/1500 [00:00<?, ? examples/s]
Map: 0% | 0/750 [00:00<?, ? examples/s]
1 हमने उसका जन्मदिन मनाया
2 साउथ दिल्ली नगर निगम सख्त शॉपिंग मॉल के बाहर नहीं देखेंगे होर्डिंग
3 उत्तर कोरिया ने अमेरिका को दी हमले की धमकी
4 अगले कमरे में अनेक रोमन मूर्तियाँ हैं
5 तुम ने टॉम को कहाँ भेज दिया
6 सर्दियों के आने से दिन छोटे होते जाते हैं
7 मुझे और वक्रत दो
8 कंगना के वकील ने कहा पुलिस ने किसी लैपटॉप की डिमांड नहीं की है
9 क्या सवाल है
10 वह अच्छा राजा था

```

```
import os
import json
```

```
def extract_all_chars(batch):
    all_text = " ".join(batch["sentence"])
    return {"all text": [all_text]}
```

```
vocabs = common_voice_train.map(
    extract_all_chars, batched=True, batch_size=-1, remove_columns=common_voice_train.column_names
)
all_text = " ".join(vocabs["all_text"])
vocab_list = sorted(list(set(list(all_text))))
```

```
# Remove the space from the set; we'll add a dedicated word_delimiter_token later.
if " " in vocab_list:
    vocab_list.remove(" ")

# Remove English letters (latin script)
vocab_list = [c for c in vocab_list if not (c >= 'a' and c <= 'z') and not (c >= 'A' and c <= 'Z')]
```

```
# Build vocab dict
vocab_dict = {v: k for k, v in enumerate(vocab_list)}
vocab_dict["|"] = len(vocab_dict) # word delimiter
vocab_dict["[UNK]"] = len(vocab_dict)
vocab_dict["[PAD]"] = len(vocab_dict)

print("Vocab size:", len(vocab_dict))
print("Sample of vocab keys:", list(vocab_dict.keys())[60])
```

```
# Save vocab to disk
os.makedirs(OUTPUT_DIR, exist_ok=True)
vocab_path = os.path.join(OUTPUT_DIR, "vocab.json")
with open(vocab_path, "w", encoding="utf-8") as f:
    json.dump(vocab_dict, f, ensure_ascii=False, indent=2)
print("Saved vocab to:", vocab_path)
```

Map: 0% | 0/1500 [00:00<?, ? examples/s]
 Vocab size: 73
 Sample of vocab keys: ['ँ', 'ं', 'ः', 'अ', 'आ', 'इ', 'ई', 'उ', 'ऊ', 'ऋ', 'ए', 'ऐ', 'ऑ', 'औ', 'औ', 'क', 'ख', 'ग', 'घ', 'च', ' ']
 Saved vocab to: Ed-168/wav2vec2-large-xls-r-300m-hi\vocab.json

```
from transformers import Wav2Vec2CTCTokenizer
tokenizer = Wav2Vec2CTCTokenizer.from_pretrained(
    OUTPUT_DIR,
    unk_token="[UNK]",
    pad_token="[PAD]",
    word_delimiter_token="|"
)

from transformers import Wav2Vec2FeatureExtractor

feature_extractor = Wav2Vec2FeatureExtractor(feature_size=1, sampling_rate=16000, padding_value=0.0, do_normalize=True, return_attention_weights=True)

from transformers import Wav2Vec2Processor

processor = Wav2Vec2Processor(feature_extractor=feature_extractor, tokenizer=tokenizer)
```

```
common_voice_train = common_voice_train.cast_column("audio", Audio(sampling_rate=16_000))
common_voice_test = common_voice_test.cast_column("audio", Audio(sampling_rate=16_000))
print(common_voice_train[0]["audio"])
```

{'path': 'C:\\Users\\EDWIN\\cache\\huggingface\\datasets\\downloads\\extracted\\2c6cd998a8800b56f2fb15d7259927dbca1bd0fa2f05133309c-1.31181878e-07, 2.62807589e-07, 4.76284185e-08]], 'sampling_rate': 16000}

```
import IPython.display as ipd
import numpy as np
import random

rand_int = random.randint(0, len(common_voice_train)-1)

print(common_voice_train[rand_int]["sentence"])
ipd.Audio(data=common_voice_train[rand_int]["audio"]["array"], autoplay=True, rate=16000)
```

दूसरा मकान खरीदने पर कितना मिलता है इनकम टैक्स में फायदा

0:00 / 0:05

```
rand_int = random.randint(0, len(common_voice_train)-1)

print("Target text:", common_voice_train[rand_int]["sentence"])
print("Input array shape:", common_voice_train[rand_int]["audio"]["array"].shape)
print("Sampling rate:", common_voice_train[rand_int]["audio"]["sampling_rate"])
```

Target text: कहाँ गई हो
 Input array shape: (42048,)
 Sampling rate: 16000

```
def prepare_dataset(batch):
    audio = batch["audio"]

    # batched output is "un-batched"
    batch["input_values"] = processor(audio["array"], sampling_rate=audio["sampling_rate"]).input_values[0]
    batch["input_length"] = len(batch["input_values"])

    with processor.as_target_processor():
        batch["labels"] = processor(batch["sentence"]).input_ids
    return batch

common_voice_train = common_voice_train.map(prepare_dataset, remove_columns=common_voice_train.column_names)
common_voice_test = common_voice_test.map(prepare_dataset, remove_columns=common_voice_test.column_names)
```

```

Map: 0%|          | 0/1500 [00:00<?, ? examples/s]
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.venv\lib\site-packages\transformers\models\wav2vec2\processing_wav2vec2.py
warnings.warn(
Map: 0%|          | 0/750 [00:00<?, ? examples/s]

# max_input_length_in_sec = 5.0
# common_voice_train = common_voice_train.filter(lambda x: x < max_input_length_in_sec * processor.feature_extractor.sampling_rate, inpu

len(common_voice_train)

1500

# import torch

# from dataclasses import dataclass, field
# from typing import Any, Dict, List, Optional, Union

@dataclass
class DataCollatorCTCWithPadding:
    processor: Wav2Vec2Processor
    padding: Union[bool, str] = True

    def __call__(self, features: List[Dict[str, Union[List[int], torch.Tensor]]]) -> Dict[str, torch.Tensor]:
        # split inputs and labels since they have to be of different lengths and need
        # different padding methods
        input_features = [{"input_values": feature["input_values"]} for feature in features]
        label_features = [{"input_ids": feature["labels"]} for feature in features]

        batch = self.processor.pad(
            input_features,
            padding=self.padding,
            return_tensors="pt",
        )
        with self.processor.as_target_processor():
            labels_batch = self.processor.pad(
                label_features,
                padding=self.padding,
                return_tensors="pt",
            )

        # replace padding with -100 to ignore loss correctly
        labels = labels_batch["input_ids"].masked_fill(labels_batch.attention_mask.ne(1), -100)

        batch["labels"] = labels

        return batch

data_collator = DataCollatorCTCWithPadding(processor=processor, padding=True)

wer_metric = evaluate.load("wer")

def compute_metrics(pred):
    # pred.predictions is float logits of shape (batch, time, vocab_size)
    pred_logits = pred.predictions
    pred_ids = torch.from_numpy(pred_logits).argmax(-1)

    # Decode predictions and references
    pred_str = processor.batch_decode(pred_ids, skip_special_tokens=True)
    # Replace -100 with pad_token_id for decoding refs
    label_ids = pred.label_ids
    label_ids[label_ids == -100] = processor.tokenizer.pad_token_id
    label_str = processor.batch_decode(label_ids, group_tokens=False)

    wer = wer_metric.compute(predictions=pred_str, references=label_str)
    return {"wer": wer}

from transformers import Wav2Vec2ForCTC

model = Wav2Vec2ForCTC.from_pretrained(
    "facebook/wav2vec2-xls-r-300m",
    attention_dropout=0.0,
    hidden_dropout=0.0,
    feat_proj_dropout=0.0,
    mask_time_prob=0.05,

```

```

        layerdrop=0.0,
        ctc_loss_reduction="mean",
        pad_token_id=processor.tokenizer.pad_token_id,
        vocab_size=len(processor.tokenizer),
    )

```

Some weights of Wav2Vec2ForCTC were not initialized from the model checkpoint at facebook/wav2vec2-xls-r-300m and are newly initialized. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
model.freeze_feature_extractor()
```

c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.venv\lib\site-packages\transformers\models\wav2vec2\modeling_wav2vec2.py
warnings.warn(

```
repo_name = "Ed-168/wav2vec2-large-xls-r-300m-hi"
```

```
from transformers import TrainingArguments
```

```

training_args = TrainingArguments(
    output_dir=repo_name,
    group_by_length=True,
    per_device_train_batch_size=16,
    gradient_accumulation_steps=2,
    num_train_epochs=40,
    gradient_checkpointing=True,
    fp16=True,
    save_steps=400,
    eval_steps=400,
    logging_steps=50,
    learning_rate=3e-4,
    warmup_steps=500,
    save_total_limit=2,
    push_to_hub=False,
)

```

```
from transformers import Trainer
```

```

trainer = Trainer(
    model=model,
    data_collator=data_collator,
    args=training_args,
    compute_metrics=compute_metrics,
    train_dataset=common_voice_train,
    eval_dataset=common_voice_test,
    tokenizer=processor.feature_extractor,
)

```

C:\Users\EDWIN\AppData\Local\Temp\ipykernel_27200\2066899627.py:3: FutureWarning: `tokenizer` is deprecated and will be removed in a future version of Transformers.
trainer = Trainer(

```
trainer.train()
```

[1880/1880 12:00:35, Epoch 40/40]

```
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\transformers\models\wav2vec2\processing_wav2vec2.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\torch\utils\checkpoint.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\transformers\models\wav2vec2\processing_wav2vec2.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\torch\utils\checkpoint.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\transformers\models\wav2vec2\processing_wav2vec2.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\torch\utils\checkpoint.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\transformers\models\wav2vec2\processing_wav2vec2.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.env\lib\site-packages\torch\utils\checkpoint.py:85: UserWarning: None of the vocab files were found in the cache. Falling back to use the local files.
warnings.warn(
TrainOutput(global_step=1880, training_loss=1.4051658075540623, metrics={'train_runtime': 43248.0009, 'train_samples_per_second': 1.387, 'train steps per second': 0.043, 'total flops': 9.181862684461875e+18, 'train loss': 1.4051658075540623, 'epoch': 40.0})
```



```

batch["input_length"] = len(batch["input_values"])

with processor.as_target_processor():
    batch["labels"] = processor(batch["sentence"]).input_ids
return batch

#common_voice_train = common_voice_train.map(prepare_dataset, remove_columns=common_voice_train.column_names)
common_voice_test = common_voice_test.map(prepare_dataset, remove_columns=common_voice_test.column_names)

Map: 0%|          | 0/750 [00:00<?, ? examples/s]
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.venv\lib\site-packages\transformers\models\wav2vec2\processing_wav2vec2.py:143: UserWarning:
  warnings.warn(

from datasets import load_dataset
import soundfile as sf
import torch

# Load only a small sample from the Common Voice Hindi test set
# common_voice_test = load_dataset("mozilla-foundation/common_voice_16_0", "hi", split="test")

input_dict = processor(common_voice_test[1]["input_values"], return_tensors="pt", padding=True)

logits = model(input_dict.input_values.to("cuda")).logits

pred_ids = torch.argmax(logits, dim=-1)[0]

It is strongly recommended to pass the `sampling_rate` argument to `Wav2Vec2FeatureExtractor()`. Failing to do so can result in silent predictions.

print("Prediction:")
print(processor.decode(pred_ids))

common_voice_test_prediction = load_dataset("mozilla-foundation/common_voice_16_0", "hi", split="test")
print("\nReference:")
print(common_voice_test_prediction[1]["sentence"])

Prediction:
आबढ़ामपूढ़ में अखिलेश बाटेंगे लैपटॉक का लॉलीपॉट

Reference:
अब रामपुर में अखिलेश बाटेंगे लैपटॉप का 'लॉलीपॉप'

repo_name = "Ed-168/wav2vec2-large-xls-r-300m-hi"
tokenizer.push_to_hub(repo_name)
tokenizer.save_pretrained(repo_name)

README.md: 0.00B [00:00, ?B/s]
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.venv\lib\site-packages\huggingface_hub\file_download.py:143: UserWarning:
  warnings.warn(message)
('Ed-168/wav2vec2-large-xls-r-300m-hi\tokenizer_config.json',
 'Ed-168/wav2vec2-large-xls-r-300m-hi\special_tokens_map.json',
 'Ed-168/wav2vec2-large-xls-r-300m-hi\vocab.json',
 'Ed-168/wav2vec2-large-xls-r-300m-hi\added_tokens.json')

repo_name = "Ed-168/wav2vec2-large-xls-r-300m-hi"
model.push_to_hub(repo_name)
processor.push_to_hub(repo_name)

Processing Files (0 / 0)           : |           | 0.00B / 0.00B
New Data Upload                   : |           | 0.00B / 0.00B
...Temp\tmpgb9j8hra\model.safetensors: 4%|3      | 50.4MB / 1.26GB
README.md: 0.00B [00:00, ?B/s]
c:\Users\EDWIN\OneDrive\Documents\GitHub\Multilingual-ASR\.venv\lib\site-packages\huggingface_hub\file_download.py:143: UserWarning:
  warnings.warn(message)

```